

Classification of Histology and Pathology Images Using Convolutional Neural Networks

Luong Nguyen
University of Pittsburgh
Biomedical Building Tower 3
Fifth Ave
Pittsburgh, PA 15026
luongn@andrew.cmu.edu

Dan Spagnolo
University of Pittsburgh
Biomedical Building Tower 3
Fifth Ave
Pittsburgh, PA 15026

Jakob Bauer
Carnegie Mellon University
5000 Forbes Avenue
Pittsburgh, PA 15213
jsbauer@andrew.cmu.edu

1. Introction

Deep learning has accomplished multiple successes with natural images, including segmenting, tagging, and scene understanding. However, only recently has deep learning ventured into the realm of medical images with multiple papers appeared in the top-tiered conference in medical imaging [2, 4]. The slower adaptation of deep learning in medical image analysis could be explained by several limitations of this field, including expensive image capture, cost prohibitive annotations, and the lack of publicly available datasets with annotations. It takes eight years to train a pathologist and one year is spent on learning histology (i.e., recognizing images of normal tissues from different organs). This motivates us to build an automated system to classify tissues from different organs.

2. Aims

Facing these challenges, our team sets out to answer the following questions on learning from these images: (1) Can we train a convolutional neural network (CNN) to differentiate normal tissues from different organs? If so, what are the differentiating features of tissues? (2) Can we train a CNN to differentiate between normal and cancerous tissues from the same organs? If so, what are the differentiating features? (3) Can we train a CNN to differentiate cancerous tissues from different organs? If so, what are the differentiating features? (4) can we improve our classifier to a significant degree using augmentation methods to increase our training data? If our CNN could not achieve any of the above goals, even with data augmentation, where did we go wrong?

3. Dataset

We curated a set of 69 whole slide images (WSIs) of tissues from 25 different organs. They are microscope images

of hematoxylin and eosin (H&E stained) tissues in which H stains nuclei purple and E stains cytoplasm and connective tissue pink. The WSIs were collected from UPMC Shady-side and scanned using Aperio XT scanner. Each WSI is of size 20,000 x 30,000 pixels at 20x magnification. These images will be divided into overlapping patches of size 256 x 256 pixels and only patches with tissues (non-empty space) are included in the analysis, resulting in about 4000 patches/WSI. In addition, for each organ, we will download 2 WSIs from the publicly available Cancer Genome Atlas (TCGA) to answer the questions related to cancerous tissues.

4. Methods

We will try multiple methods for data augmentations to increase the number of patches per WSI up to 10000 patches. These methods may include rotating and flipping patches, scale jittering, modifying color and brightness, as well as adding noise to patches [3]. In addition to ground truth, we will collect non-expert naive observer (Jakob)s and intermediate observer (Daniel)s annotations for comparison with human classification. For baseline comparison with shallow learning methods, we can train an SVM with SIFT/HOG features [1]. We will experiment with standard CNN architectures such as AlexNet, VGGNet, and LeNet. We suspect that our problem will not require as many CNN layers as 1000-class datasets, for example, and thus we will experiment with CNN depth. We will use Torch or Caffe to train our CNNs. The last layer will be removed to generate feature vectors. These feature vectors will serve as input to a classifier (SVM, logistic regression) to determine the tissue origin and disease status. Finally we will conduct error analysis to find out where we can improve our CNN. The metrics for comparisons will be precision-recall curves and F1 scores. Our goal is to outperform our shallow learning baseline as well as our naive and intermediate observers.

5. Expected Complications

From our experience with classes in histology, it is possible to train a human to accomplish aim (1) fairly well after a one-semester class (Luong and Dan took a histology class). The ability to differentiate normal tissues from different organs is based on recognizing landmark features of different organs. For examples, kidney tissues have outer and inner cortex while breast tissues have clusters of milk conducting ducts (see Figure 1, left). Therefore, we are convinced that it is plausible to teach computers to read normal histology images. However, patches taken from WSI do not always contain landmark features of organs and tissue types. For example, adipose (fat) tissue is the same for many organs. We may not be able to properly classify image patches if these landmark features are not sampled in a given patch. This implies high false positive rate if we try to classify patches into different tissues. In addition, we might encounter difficulties analyzing cancerous tissues since invasive cancer can completely destroy the unique structures of different tissue types. Therefore, we expect to have a hard time accomplishing the task in aim (3) since destroyed tissue structures can look very similar across different organs. Another complication can come from the color variation of these tissue images since different hospitals and pathologists have different slide preparation procedures. This means the CNN can focus on figuring out different color schemes for classification instead of actually discovering different structures of tissues. To resolve this, we will apply some standard color normalization for histology images.

Finally, there is some concern for the number of training and testing images we will have. Due to the cost of data acquisition in this domain, our dataset scale is much smaller than would be found in typical datasets in the computer vision field. We will attempt to overcome this difficulty using the data augmentation techniques described in Methods. If our shallow learning methods outperform the deep learning approach, this may be indicative of insufficient data.




Figure 1. Breast tissue with ducts and lobules.




Figure 2. Bone tissue with large area of pink stained calcium.

References

- [1] H. Chang, A. Borowsky, P. Spellman, and B. Parvin. Classification of tumor histology via morphometric context. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2203–2210, 2013.
- [2] A. A. Cruz-Roa, J. E. A. Ovalle, A. Madabhushi, and F. A. G. Osorio. A deep learning architecture for image representation, visual interpretability and automated basal-cell carcinoma cancer detection. In *Medical Image Computing and Computer-Assisted Intervention–MICCAI 2013*, pages 403–410. Springer, 2013.
- [3] D. Eigen and R. Fergus. Predicting depth, surface normals and semantic labels with a common multi-scale convolutional architecture. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 2650–2658, 2015.
- [4] H. R. Roth, L. Lu, A. Farag, H.-C. Shin, J. Liu, E. B. Turkbey, and R. M. Summers. Deeporgan: Multi-level deep convolutional networks for automated pancreas segmentation. In *Medical Image Computing and Computer-Assisted Intervention–MICCAI 2015*, pages 556–564. Springer, 2015.